Progressive Strategies for Monte-Carlo Tree Search

5 authors, including:

Guillaume Chaslot
Maastricht University
32 PUBLICATIONS   2,141 CITATIONS

SEE PROFILE

Mark H.M. Winands
Maastricht University
158 PUBLICATIONS   2,660 CITATIONS

SEE PROFILE

H. Jaap Van Den Herik
Leiden University
627 PUBLICATIONS   7,374 CITATIONS

SEE PROFILE

Jos W. H. M. Uiterwijk
Maastricht University
178 PUBLICATIONS   2,799 CITATIONS

SEE PROFILE

PROGRESSIVE STRATEGIES
FOR MONTE-CARLO TREE SEARCH

GUILLAUME M.J-B. CHASLOT, MARK H.M. WINANDS,

H. JAAP VAN DEN HERIK and JOS W.H.M. UITERWIJK

MICC-IKAT, Games and AI Group, Faculty of Humanities and Sciences,
Universiteit Maastricht, P.O. Box 616, 6200 MD Maastricht, The Netherlands
{g.chaslot, m.winands, herik, uiterwijk}@micc.unimaas.nl

BRUNO BOUZY

Centre de Recherche en Informatique de Paris 5, Université Paris 5 Descartes,
45, rue des Saints Pères, 75270 Cedex 06, France

bouzy@math-info.univ-paris5.fr

Monte-Carlo Tree Search (MCTS) is a new best-first search guided by the results of Monte-Carlo simulations. In this article we introduce two *progressive strategies* for MCTS, called progressive bias and progressive unpruning. They enable the use of relatively time-expensive heuristic knowledge without speed reduction. Progressive bias directs the search according to heuristic knowledge. Progressive unpruning first reduces the branching factor, and then increases it gradually again. Experiments assess that the two progressive strategies significantly improve the level of our Go program Mango. Moreover, we see that the combination of both strategies performs even better on larger board sizes.

1. Introduction

Over fifty years, two-person zero-sum games with perfect information have been addressed by many AI researchers with great success.[17] The classical approach is to use the alpha-beta framework, combined with a positional evaluation function. Such an evaluation function is applied to the leaf nodes of a search tree. If the node represents a terminal position (or an endgame database position) it produces an exact value. Otherwise, heuristic knowledge is used to estimate the value of the leaf node. This technique led to excellent results in many games (e.g., chess and checkers).[7,21,20]

However, in several games building an evaluation function based on heuristic knowledge for a non-terminal position is a difficult and time-consuming issue; the most notorious example is the game of Go.[2] It is probably one of the reasons why Go programs so far did not achieve a strong level, despite intensive research and additional use of knowledge-based methods.

Recently, researchers proposed to use Monte-Carlo simulations as an evaluation function.[5,6] Yet, this approach remained too slow to achieve a satisfying search depth. Even more recently, three slightly different uses of Monte-Carlo simulations

within a tree-search context have been proposed.[10,13,18] The new general method, which we call "Monte-Carlo Tree Search" (MCTS), resulted from it. MCTS is not a classical tree search followed by a Monte-Carlo evaluation, but rather a best-first search guided by the results of Monte-Carlo simulations. This method uses two main strategies, which aim at different purposes described below. (1) A *selection strategy*, derived from the Multi-Armed Bandit problem, is able to increase the quality of the chosen moves in the tree when the number of simulations grows.[11,15] Yet, the strategy requires the results of several previous simulations. (2) When a sufficient amount of results is not available, a *simulation strategy* decides on the moves to be played.[1,15] So, the challenge is: how do we harmonize a simulation strategy (necessary to be applied by a lack of sufficient results) with the selection strategy?

In this article, we propose two progressive strategies as a soft transition between the simulation strategy and the selection strategy. The strategies enable, among others, the use of time-consuming heuristic knowledge. Below, we use the game of Go as test domain. Go is challenging, because so far programs are not able to defeat expert humans, and thus it has been a testbed for artificial-intelligence techniques for over 30 years.

The article is organized as follows. In Section 2, we present the Monte-Carlo Tree Search method. In Section 3, we describe the two progressive strategies. Section 4 presents the experiments, performed with the Go program Mango. Section 5 discuss alternative solutions. Section 6 summarizes the contributions, formulates conclusions, and gives an outlook on future research.

2. Monte-Carlo Tree Search

Monte-Carlo Tree Search (MCTS) is a best-first search method, which does not require a positional evaluation function. It is based on a randomized exploration of the search space: in the beginning of the search, exploration is performed fully at random. Then, using the results of previous explorations, the algorithm becomes able to predict the most promising moves more accurately, and thus, their evaluation becomes more precise. The basic structure of MCTS is given in Subsection 2.1. Relevant pseudo-code is provided in Subsection 2.2. Four strategic tasks are discussed in Subsection 2.3. Finally, we discuss how to select the move to be played in the actual game in Subsection 2.4.

2.1. *Structure of MCTS*

In MCTS, each node i represents a given position (also called a state) of the game. A node contains at least the following two pieces of information: (1) the current value v_i of the position (usually the average of the results of the simulated games that visited this node), and (2) the visit count of this position n_i. MCTS usually starts with a tree containing only the root node.
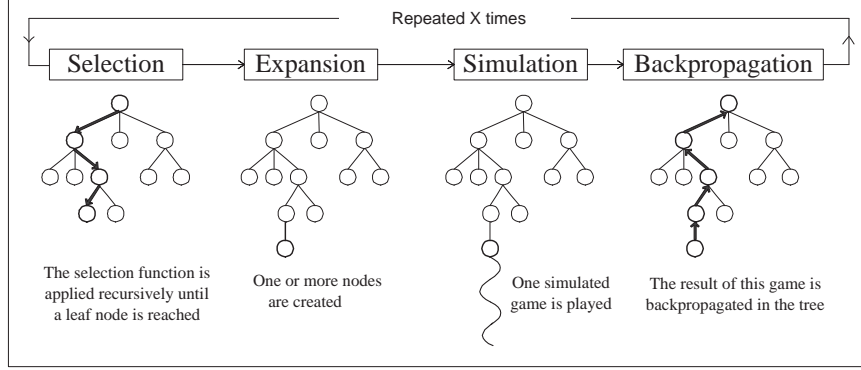
Fig. 1. Outline of a Monte-Carlo Tree Search.

MCTS consists of four steps, repeated as long as there is time left. The steps are as follows. (1) The tree is traversed from the root node to a leaf node (L), using a *selection strategy*. (2) An *expansion strategy* is called to store one (or more) children of L in the tree. (3) A *simulation strategy* plays moves in self-play until the end of the game is reached. The result R of this "simulated" game is $+1$ in case of a win for Black (the first player in Go), 0 in case of a draw, and -1 in case of a win for White. (4) R is propagated back through the tree according to a *backpropagation strategy*. Finally, the move played by the program is the child of the root with the highest visit count. The four steps of MCTS are explained in some detail in Figure 1, and more elaborated in Subsection 2.3.

2.2. *Relevant pseudo-code*

The pseudo-code for MCTS is given in Figure 2. In this algorithm, $\mathcal{ST}$ is the set of all nodes (the search tree), *Select(Node N)* is the selection function, which returns one child of the node N. *Expand(Node N)* is the function that stores one more node in the tree, and returns this node. *Play_simulated_game(Node N)* is the function that plays a simulated game from the node N, and returns the result $R \in \{-1, 0, 1\}$ of this game. *Backpropagate(Integer R)* is the procedure that updates the value of the node depending on the result R of the last simulated game. $\mathcal{N}_c$*(node N)* is the set of the children of the node N.

2.3. *The four strategic tasks*

As has been mentioned earlier, the four strategic tasks in MCTS are selection, expansion, simulation, and backpropagation. They are each discussed in detail below. For each, we will show how we use them in our Go program MANGO.

4   *G.M.J-B. Chaslot, M.H.M. Winands, H.J. van den Herik, J.W.H.M. Uiterwijk, and B. Bouzy*

void MCTS(Node root_node)

```
1  while(has_time)
2  {
3        current_node ← root_node
4        while (current_node ∈ ST)
5        {
6              last_node ← current_node
7              current_node ← Select(current_node)                    // Selection
8        }
9        last_node ← Expand(last_node)                                // Expansion
10       R ← Play_simulated_game(last_node)                          // Simulation
11       while(current_node ∈ ST)
12       {
13             current_node.Backpropagate(R)                         // Backpropagation
14             current_node.visit_count ← current_node.visit_count + 1
15             current_node ← current_node.parent
16       }
17 }
18 return best_move = argmax_{N∈𝒩_c(root_node)}(N.visit_count)
```

Fig. 2. Pseudo-code for Monte-Carlo Tree Search.

2.3.1. *Selection*

Selection is the strategic task that selects one of the children of a given node. It controls the balance between exploitation and exploration. We explain both notions below. Exploitation is the task to select the move that leads to the best results so far. Exploration deals with less promising moves that still have to be examined, due to the uncertainty of the evaluation. Similar balancing of exploitation and exploration has been studied in the literature, in particular with respect to the Multi-Armed Bandit (MAB) problem.[19] The MAB problem considers a gambling device and a player, whose objective is maximizing the reward from the device. At each time step, the player can select one of N arms of the device, which gives a reward. In most settings, the reward obeys a stochastic distribution. The selection problem of MCTS could be viewed as a MAB problem for a given node: the problem is to select the next move (arm) to play, which will give an unpredictable reward (the outcome of a single random game). Knowing the past results, the problem is to find the optimal move. However, the main difference with the MAB problem is that MCTS works by using sequentially several selections: the selection at the root node, the selection at depth one, the selection at depth two, etc. Several algorithms have been designed for this setup,[10,11,13] or have been derived from MAB algorithms.[15,18]

Selection strategy used in Mango
We use the strategy UCT (Upper Confidence Bound applied to Trees).[18] This strategy is easy to implement, and used in many programs. UCT works as follows. Let I be the set of nodes reachable from the current node p. UCT selects the child k of

the node p that satisfies formula 2.1:

$$k \in argmax_{i \in I}\left(v_i + C \times \sqrt{\frac{\ln n_p}{n_i}}\right) \tag{2.1}$$

where v_i is the value of the node i, n_i is the visit count of i, and n_p is the visit count of p. C is a coefficient, which has to be tuned experimentally. In Mango, we use $C = 0.7$. *In practice, UCT is only applied in nodes of which the visit count is higher than a certain threshold T* (i.e., 30 in Mango).[13] If the node has been visited fewer times than this threshold, the next node is selected according to the *simulation strategy*, discussed in 2.3.3.

2.3.2. *Expansion*

Expansion is the strategic task that, for a given leaf node L, decides whether this node will be expanded by storing one or more of its children in memory. The simplest rule is to expand one node per simulated game.[13] The expanded node corresponds to the first position encountered that was not stored yet. Note that the expansion could also be performed after the simulation.

Expansion strategy used in Mango
In addition to expanding one node per simulated game, we also expand all the children of a node when a node's visit count equals T.

2.3.3. *Simulation*

Simulation (also called *playout*) is the strategic task that selects moves in self-play until the end of the game. This task might consist of playing plain random moves or – better – pseudo-random moves chosen according to a *simulation strategy*. Indeed, the use of an adequate simulation strategy has been shown to improve the level of play significantly.[1,16] The main idea is to play interesting moves by using patterns, capture considerations, and proximity to the last move. The simulation requires that the number of moves per game is limited. When considering the game of Go, extra rules are added to satisfy this condition: (1) a player should not play in his eyes, and (2) the game is stopped and scored if it exceeds a given number of moves. Elaborating an efficient simulation strategy is a difficult issue. If the strategy is too stochastic (e.g., if it selects moves nearly randomly), then the moves played are often weak, and the level of the Monte-Carlo program is decreasing. In contrast, if the strategy is too deterministic (e.g., if the selected move for a given position is almost always the same, i.e., too much exploitation takes place) then the exploration of the search space becomes too selective, and the level of the Monte-Carlo program is decreasing too.

6   *G.M.J-B. Chaslot, M.H.M. Winands, H.J. van den Herik, J.W.H.M. Uiterwijk, and B. Bouzy*

Simulation strategy used in MANGO

Let $\mathcal{M}$ be the set of all possible moves for a given position. Each move $j \in \mathcal{M}$ is given an urgency $U_j \geq 1$. The simulation strategy selects one move from $\mathcal{M}$. The probability of each move to be selected is $p_j = \frac{U_j}{\sum_{k \in \mathcal{M}} U_k}$. The urgency is the sum of two values: the capture value and the pattern value.

(1) *Capture value.* This value is given to moves capturing stones or preventing captures. It equals $50,000 \times$ the number of captured stones, plus $5,000 \times$ the number of stones prevented from capture. Using a capture value was first proposed by Bouzy, and later improved successively by Coulom and Cazenave.[1,13,8]

(2) *Pattern value.* For each possible 3×3 pattern, the value of the central move has been learned by a dedicated algorithm developed in our previous research.[4] The pattern values range from 1 to 2433.

Moves within a Manhattan distance of 1 from the previous move have their urgency multiplied by 20. This idea is taken from the strategy developed by Gelly and Wang[16] and slightly adapted.

2.3.4. *Backpropagation*

Backpropagation is the procedure that propagates the *result* of a simulated game k backwards from leaf node L to the nodes it had to traverse to reach this leaf node. This result is counted positively ($R_k = +1$) if the game is won, and negatively ($R_k = -1$) if the game is lost. Draws lead to a result $R_k = 0$. The *value* v_L of a node is computed by taking the average of the results of all simulated games made through this node, i.e., $v_L = (\sum_k R_k)/n_L$. Several backpropagation strategies have been proposed in the literature.[9,11,13] However, the best results in game play have been obtained by using the plain average of the simulations.

Backpropagation strategy used in MANGO

In MANGO we use the plain average strategy described above.

2.4. *Final move selection*

After the simulations, the move finally played by the program in the actual game is the one corresponding to the "best child" of the root. There are different ways to define which child is the best.

(1) *Max child.* The max child is the child that has the highest value.

(2) *Robust child.* The robust child is the child with the highest visit count.

(3) *Robust-max child.* The robust-max child is the child with both the highest visit count and the highest value. If there is no robust-max child at the moment, more simulations are played until a robust-max child occurs.[13]

(4) *Secure child.* The secure child is the child that maximizes a lower confidence bound, i.e., which maximizes the quantity $v + \frac{A}{\sqrt{n}}$, where A is a parameter (set to 4 in our experiments), v is the node's value, and n is the node's visit count.

Final move selection used in MANGO

MANGO uses the robust child. In our preliminary experiments we did not measure a significant difference between the methods discussed above, when a sufficient amount of simulations per move was played. However, when only a short thinking time per move was used (e.g., below a second), choosing the max child turned out to be significantly weaker than other methods.

3. Progressive Strategies

When a node has been visited only a few times, a well-tuned simulation strategy chooses moves more accurately than a selection strategy. However, when a node has been visited quite often, the selection strategy is more accurate.[10,11,13,18]

We propose a "progressive strategy" that performs a soft transition between the simulation strategy and the selection strategy. Such a strategy uses (1) the information available for the selection strategy, and (2) some time-expensive domain knowledge. A progressive strategy is similar to a simulation strategy when a few games have been played, and converges to a selection strategy when numerous games have been played.

In the following two subsections we describe the two progressive strategies used in our Go-playing program MANGO: *progressive bias* (Subsection 3.1) and *progressive unpruning* (Subsection 3.2). Subsection 3.3 describes the heuristic domain knowledge used in MANGO. Subsection 3.4 discusses the time efficiency of these heuristics.

3.1. *Progressive bias*

The aim of the *progressive bias* strategy is to direct the search according to – possibly time-expensive – heuristic knowledge. For that purpose, the selection strategy is modified according to that knowledge. The influence of this modification is important when a few games have been played, but decreases fast (when more games have been played) to ensure that the strategy converges to a selection strategy. We modified the UCT selection in the following way. Instead of selecting the move which satisfies formula 2.1, we select the node k which satisfies formula 3.2. We call this formula our enhancement.

$$k \in argmax_{i \in I} \left(v_i + C \times \sqrt{\frac{\ln n_p}{n_i}} + f(n_i) \right) \qquad (3.2)$$

In MANGO, we chose $f(n_i) = \frac{H_i}{n_i + 1}$, where H_i is a coefficient representing heuristic knowledge, which depends only on the board configuration represented by the node i. The variables n_p and n_i, and coefficient C are the same as in Section 2.

More details on the construction of H_i are given in Subsection 3.3. Formula 3.2 has the following four properties.

(1) When the number of games n_p made through a node p equals T (i.e., 30 in MANGO), the selection algorithm starts to be applied in this node. For all the children i of this node with $n_i = 0$, $\sqrt{\frac{\ln n_p}{n_i}}$ is replaced by a fixed number M (e.g., 100 in MANGO) satisfying $\forall i, M \gg v_i$. v_i is replaced by 0 when $n_i = 0$. Thereafter, the algorithm selects every unexplored child once. The order in which these children are selected is given by $f(n_i)$, i.e., the children with the highest heuristic values are selected first.
(2) If only a few simulations have been made through the node (e.g., from around 30 to 100 in MANGO), and if the heuristic value H_i is sufficiently high, the term $\frac{H_i}{n_i+1}$ is dominant. Hence, the number of simulations made depends more on the domain knowledge H_i than on the results of the simulated games. It is an advantage to use mainly the domain knowledge at this stage, because then the results of only a few simulated games are affected by a large uncertainty. The behavior of the algorithm is therefore close to the behavior of a simulation strategy.
(3) When the number of simulations increases (e.g., from around 100 to 500 in MANGO), both the results of the simulated games and the domain knowledge have a balanced impact on the selection.
(4) When the number of simulations is high (e.g., > 500 in MANGO), the influence of the domain knowledge is low compared to the influence of the previous simulations, because the domain knowledge decreases by $O(1/n_i)$, and the term corresponding to the simulation decreases by $O(\sqrt{\ln n_p/n_i})$. The behavior of the algorithm is, at this point, close to the behavior of a classical selection strategy (UCT). The only difference with plain UCT occurs if two positions i and j have the same value $v_i = v_j$, but different heuristic evaluations H_i and H_j. Then, the position with the highest heuristic evaluation will be selected more often.

3.2. *Progressive unpruning*

We have seen in MANGO that when there is not much time available and simultaneously the branching factor is high, MCTS performs poorly. Our solution, *progressive unpruning*, consists of (1) reducing the branching factor artificially when the selection function is applied, and (2) increasing it progressively as more time becomes available. When the number of games n_p in a node p equals the threshold T, progressive unpruning "prunes"[a] most of the children. The children, which are not pruned from the beginning, are the k_{init} children with the highest heuristic values. In MANGO k_{init} was set to 5. The children of the node i are progressively "un-

[a]A node is pruned if it cannot be accessed in the simulated games.

pruned". In MANGO, k nodes are unpruned when the number of simulations in the parent surpasses $A \times B^{k-k_{init}}$ simulated games. A was set experimentally to 50 and B to 1.3. The scheme is shown in Figure 3.


Fig. 3. Progressive unpruning in MANGO.

3.3. *Heuristic knowledge used in Mango*

The two previous soft-transition strategies require computing a heuristic value H_i for a given board configuration representing the node i. In this subsection we describe the heuristic, which is based on the same ideas as seen in 2.3.3. However, the heuristic knowledge for H_i is much more elaborated than the one used for the urgency value U_i. In MANGO, H_i is composed of three elements: (i) a pattern value, (ii) a capture value, and (iii) the proximity to the last moves.

The *capture value* of each move depends on (1) the number of stones that could be captured by playing the move, and on (2) the number of stones that could escape a capture by playing the move. It is calculated the same way than for the simulation strategy.

The *pattern value* is learned offline by using the pattern matching described by Bouzy and Chaslot.[3] This pattern matching was also implemented in the Go program INDIGO, and improved its level significantly.[b] In this research, each pattern assigns a value to the move that is in its center. The value corresponds to the probability that the move is played in professional games. The learning phase has been performed on $2,000$ professional games; $89,119$ patterns were learned. Each pattern contained between 0 stones (e.g., corner pattern) and 15 stones (e.g., joseki pattern). The size of the patterns was not bounded, so some patterns covered nearly the whole board, and some covered only a few intersections.

The *proximity coefficients* are computed as the Euclidean distances to all moves previously played, as shown below.

These elements are combined in the following formula to compute H_i:

$$H_i = (C_i + P_i) \sum_k \frac{1}{(2d_{k,i})^{\alpha_k}} \tag{3.3}$$

where C_i is the capture value, P_i is the pattern value, $d_{k,i}$ is the (Euclidean) distance to the k^{th} last move, and $\alpha_k = 1.25 + \frac{k}{2}$. This formula has been tuned experimentally. Computing the P_i values is the time-consuming part of the heuristic.

3.4. *Time available for heuristics*

The time consumed to compute H_i is in the order of one millisecond, which is around 1000 times slower than playing a move in a simulated game. To avoid a speed reduction in the program's performance, we compute H_i only once per node, when a certain threshold of games has been played through this node. The threshold was set to $T = 30$ in MANGO. With this setting, the speed of the program was only reduced by 4 percent. The speed reduction is low because the amount of nodes that have been visited more than 30 times is low compared to the amount of moves played in the simulated games. It can be seen in Figure 4 that the number of calls to the domain knowledge is reduced quickly as T increases. Even for $T = 9$, the number of calls to the domain knowledge is quite low compared to the number of simulated moves. The amount of nodes having a certain visit count is plotted in Figure 5. The data has been obtained from a 19×19 initial position by performing a 30-second MCTS. We have also plotted a trend line that shows that this data can be approximated by a power law.
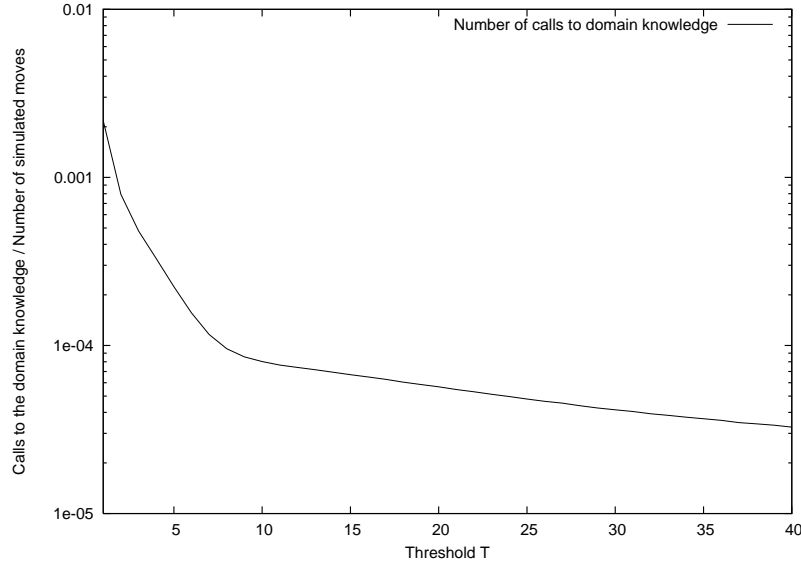

Fig. 4. Number of calls to the domain knowledge relative to the number of simulated moves, as a function of the threshold T.

Fig. 5. Number of nodes with a given visit count.

4. Experiments

Three different series of experiments were conducted. Subsection 4.1 gives the impact of each progressive strategy against GNU Go. Subsection 4.2 shows that these methods also improve the level of our program in self-play. Subsection 4.3 assesses the strength of our program Mango in recent (internet) tournaments.

4.1. Mango *vs.* GNU Go

In the first series of experiments we tested the two progressive strategies in games against GNU Go version 3.6. The experiments were performed on the 9×9, 13×13, and 19×19 boards. Our program used $20,000$ simulations per move. It takes on average less than one second on a 9×9 board, two seconds on a 13×13 board, and five seconds on a 19×19 board. The level of GNU Go has been set to 10 on the 9×9 board and on the 13×13 board, and to 0 on the 19×19 board. The results are reported in Table 1, where PB stands for progressive bias and PU for progressive unpruning.

From these experiments, the results, and our observation, we may arrive at three conclusions. First, the plain MCTS framework does not scale well to the 13×13 board and the 19×19 board, even by using GNU Go at level 0. Second, the progressive strategies increase Mango's level of play on every board size. Third, on the 19×19 board size the combination of both strategies is much stronger than each strategy applied separately.

Table 1. Results of MANGO against GNU GO.

Board size	Simulations per move	GNU Go's level	PB	PU	Wins	Games	95 percent confidence interval
9	20,000	10			33.2%	1000	3.0%
9	20,000	10		X	37.2%	1000	3.1%
9	20,000	10	X		58.3%	1000	3.1%
9	20,000	10	X	X	61.7%	2000	2.2%
13	20,000	10			8.5%	500	2.5%
13	20,000	10		X	15.6%	500	3.2%
13	20,000	10	X		30.0%	500	4.1%
13	20,000	10	X	X	35.1%	500	4.3%
19	20,000	0			0%	200	1.0%
19	20,000	0		X	3.1%	200	2.5%
19	20,000	0	X		4.8%	200	3.0%
19	20,000	0	X	X	48.2%	500	4.5%

4.2. *Self-play experiment*

We also performed self-play experiments on the different board sizes. The time setting of these experiments is 10 seconds per move. On the 9×9 board, MANGO using progressive strategies won 88 percent of 200 games played against MANGO without progressive strategies. Next, on the 13×13 board, MANGO using both progressive strategies won 81 percent of 500 games played against MANGO without progressive strategies. Finally, on the 19×19 board, MANGO using both progressive strategies won all the 300 games played against MANGO without progressive strategies. These self-play experiments show that the impact of progressive strategies is larger on the 19×19 board than on the 13×13 and 9×9 boards. This conclusion is consistent with the results of the experiments of the previous subsection.

4.3. *Tournaments*

In the last series of experiments we tested MANGO's strength by competing in computer tournaments. Table 2 presents the results by MANGO in the tournaments entered in 2007. In all these tournaments, MANGO used both progressive strategies. In this table, KGS stands for "KGS Go Server". This server is the most popular one for computer programmers, and most of the well-known programs have participated in one or more editions (e.g., MOGO, CRAZYSTONE, GO++, THE MANY FACES OF GO, GNU GO, INDIGO, AYA, DARIUSH, etc...).

As shown in the previous experiments, the progressive strategies are the main strength of MANGO. We remark that MANGO was always in the best half of the participants.

Table 2. Results by MANGO in 2007 tournaments.

Tournament	Board Size	Participants	MANGO's rank
KGS January 2007	13×13	10	2^{nd}
KGS March 2007	19×19	12	4^{th}
KGS April 2007	13×13	10	3^{rd}
KGS May 2007	13×13	7	2^{nd}
12^{th} Computer Olympiad	9×9	10	5^{th}
12^{th} Computer Olympiad	19×19	8	4^{th}
KGS July 2007	13×13	10	4^{th}

5. Discussion

An alternative enhancement to progressive bias has been proposed by Gelly and Silver.[14] It consists of introducing prior knowledge. The selected node k is the one, which satisfies formula 5.4:

$$k \in argmax_{i \in I} \left(\frac{v_i \cdot n_i + n_{prior} \cdot Q_i}{n_i + n_{prior}} + C \times \sqrt{\frac{\ln n_p}{n_i + n_{prior}}} \right) \qquad (5.4)$$

where Q_i is the prior estimation of the position. Gelly and Silver use a reinforcement learning algorithm, which learned the value from self-play on the 9×9 board.[22] n_{prior} is a coefficient that was tuned experimentally.

On the 9×9 board, this technique successfully increased MOGO's winning percentage against GNU GO from 60 percent to 69 percent. However, learning the prior value Q_i was only done for the 9×9 board. So, the scalability of this approach to larger board sizes is an open question.

We would also like to remark that a quite similar scheme to progressive unpruning, called *progressive widening*, has been proposed simultaneously by Coulom.[12] This scheme improved the level of his program CRAZYSTONE significantly. The main difference with our implementation is that the speed of unpruning as implemented in his program is slower than the one used in MANGO. This implies that the quality of the heuristic domain knowledge in his program is higher. Therefore it can afford to be more selective without pruning important moves.

6. Conclusions and Future Research

In this article we introduced the concept of progressive strategy. It enables a soft transition from a simulation to a selection strategy. Such a strategy uses (1) the information available for the selection strategy, and (2) some time-expensive domain

14   *G.M.J-B. Chaslot, M.H.M. Winands, H.J. van den Herik, J.W.H.M. Uiterwijk, and B. Bouzy*

knowledge. We have developed two progressive strategies: progressive bias and progressive unpruning. Progressive bias uses knowledge to direct the search. Progressive unpruning first reduces the branching factor, and then increases it gradually. This scheme is also dependent on knowledge.

Based on the results of the experiments performed with our program Mango, we may offer four conclusions. (1) The plain Monte-Carlo Tree Search method does not scale well to 13×13 Go, and performs even worse in 19×19 Go. (2) Progressive strategies increase the level of play of our program Mango significantly, on every board size. (3) On the 19×19 board size, the combination of both strategies is much stronger than each strategy applied separately. (4) These strategies can use relatively expensive domain knowledge with hardly any speed reduction.

For future research there are four interesting directions. First, we will investigate other progressive strategies such as *RAVE* and *UCT with prior knowledge.*[14] They have been recently proposed to include knowledge in the Monte-Carlo Tree Search framework. It would be interesting to combine them with the current progressive strategies. Second, the progressive strategies could give even better results by using more advanced knowledge, e.g., as developed by Coulom.[12] Third, another idea is to improve the heuristic knowledge by using life-and-death knowledge. Fourth, an interesting path of research is to apply the work by Coquelin and Munos[11] to improve the progressive bias.

Acknowledgments

References

1. B. Bouzy. Associating Domain-Dependent Knowledge and Monte-Carlo Approaches within a Go Program. *Information Sciences, Heuristic Search and Computer Game Playing IV*, 175(4):247–257, 2005.
2. B. Bouzy and T. Cazenave. Computer Go: An AI-oriented Survey. *Artificial Intelligence*, 132(1):39–103, 2001.
3. B. Bouzy and G.M.J-B. Chaslot. Bayesian Generation and Integration of K-nearest-neighbor Patterns for 19x19 Go. In G. Kendall and Simon Lucas, editors, *IEEE 2005 Symposium on Computational Intelligence in Games, Essex, UK*, pages 176–181, 2005.
4. B. Bouzy and G.M.J-B. Chaslot. Monte-Carlo Go Reinforcement Learning Experiments. In *IEEE 2006 Symposium on Computational Intelligence in Games, Reno, USA*, pages 187–194, 2006.
5. B. Bouzy and B. Helmstetter. Monte-Carlo Go Developments. In H.J. van den Herik, H. Iida, and E. A. Heinz, editors, *Proceedings of the 10th Advances in Computer Games Conference (ACG-10)*, pages 159–174. Kluwer Academic, 2003.
6. B. Brügmann. Monte Carlo Go. *Technical report, Physics Department, Syracuse University, http://citeseer.ist.psu.edu/655469.html*, 1993.

7. M. Campbell, A.J. Hoane Jr., and F-H. Hsu. Deep Blue. *Artificial Intelligence*, 134(1–2):57–83, 2002.

8. T. Cazenave. Playing the right atari. *ICGA Journal*, 30(1):35–42, 2007.

9. G.M.J-B. Chaslot, S. De Jong, J-T. Saito, and J.W.H.M. Uiterwijk. Monte-Carlo Tree Search in Production Management Problems. In P.-Y. Schobbens, W. Vanhoof, and G. Schwanen, editors, *Proceedings of the 18th BeNeLux Conference on Artificial Intelligence*, pages 91–98, 2006.

10. G.M.J-B. Chaslot, J-T. Saito, B. Bouzy, J.W.H.M. Uiterwijk, and H.J. van den Herik. Monte-Carlo Strategies for Computer Go. In P.-Y. Schobbens, W. Vanhoof, and G. Schwanen, editors, *Proceedings of the 18th BeNeLux Conference on Artificial Intelligence*, pages 83–90, 2006.

11. P.-A. Coquelin and R. Munos. Bandit algorithms for tree search. In *Uncertainty in Artificial Intelligence*, Vancouver, Canada, 2007.

12. R. Coulom. Computing Elo Ratings of Move Patterns in the Game of Go. In H.J. van den Herik, J.W.H.M. Uiterwijk, M.H.M. Winands, and M.P.D. Schadd, editors, *Proceedings of the Computers Games Workshop 2007 (CGW 2007)*, pages 113–124, 2007.

13. R. Coulom. Efficient selectivity and backup operators in Monte-Carlo tree search. In H.J. van den Herik, P. Ciancarini, and H.H.L.M. Donkers, editors, *Proceedings of the 5th International Conference on Computer and Games*, volume 4630 of *Lecture Notes in Computer Science (LNCS)*, pages 72–83. Springer-Verlag, Heidelberg, Germany, 2007.

14. S. Gelly and D. Silver. Combining Online and Offline Knowledge in UCT. In Zoubin Ghahramani, editor, *Proceedings of the International Conference of Machine Learning (ICML 2007)*, pages 273–280, 2007.

15. S. Gelly and Y. Wang. Exploration Exploitation in Go: UCT for Monte-Carlo Go. In *Twentieth Annual Conference on Neural Information Processing Systems (NIPS 2006)*, 2006.

16. S. Gelly, Y. Wang, R. Munos, and O. Teytaud. Modifications of UCT with Patterns in Monte-Carlo Go. Technical Report 6062, INRIA, 2006.

17. H. J. van den Herik, J. W. H. M. Uiterwijk, and J. van Rijswijck. Games Solved: Now and in the Future. *Artificial Intelligence*, 134(1–2):277–311, 2002.

18. L. Kocsis and C. Szepesvári. Bandit Based Monte-Carlo Planning. In J. Fürnkranz, T. Scheffer, and M. Spiliopoulou, editors, *Machine Learning: ECML 2006, Lecture Notes in Artificial Intelligence 4212*, pages 282–293, 2006.

19. H. Robbins. Some Aspects of the Sequential Design of Experiments. *Bulletin of the American Mathematical Society*, 55:527–535, 1952.

20. J. Schaeffer. Game Over: Black to Play and Draw in Checkers. *ICGA Journal*, 30(4):187–197, 2007.

21. J. Schaeffer, N. Burch, Y. Björnsson, A. Kishimoto, M. Müller, R. Lake, P. Lu, and S. Sutphen. Checkers is solved. *Science*, 317(5844):1518–1522, 2007.

22. D. Silver, R. S. Sutton, and M. Müller. Reinforcement Learning of Local Shape in the Game of Go. In *20th International Joint Conference on Artificial Intelligence (IJCAI 2007)*, pages 1053–1058, 2007.